

microp

GPipe

What is GPipe?

- GPipe splits large minibatches into smaller “microbatches”, and process them through model partitions (stages) running on different devices (GPUs).
-

How does it work?

划分为更小的微批次

- **Split the Minibatch:** Each input minibatch of size B is split into m equal microbatches (B/m samples per microbatch).
- **Pipelined Execution:** Process microbatches in order. A microbatch flows forward through stage 0, then stage 1, ..., until the last stage. Once a stage finishes one microbatch, it immediately starts the next.

Gradient Accumulation

- Each microbatch is trained and its gradients computed independently. Gradients from all microbatches are accumulated locally before the optimizer update—giving the same effect as if training on the full batch at once.

Limitations & Details



1. Pipeline Bubbles:

- The “fill” and “drain” phases of the pipeline (when not all devices are occupied) create idle slots or “bubbles”.
- The fraction of time lost to bubbles:

$$\text{空泡的比例} = 1 - \frac{2nm}{2n(m+n-1)} = 1 - \frac{m}{m+n-1}$$

○ m为微批的个数，n为时间片数

- Increasing the number of microbatches m (flooding the pipe), is necessary for making the bubble fraction small.




-

- Top plot is 0.75 not 0.8!

2. Memory Demand:

- Each microbatch requires storing its activations until its backward pass is finished; using more microbatches increases parallelism at the cost of memory consumption.
- In GPipe, we need to cache the activations for each microbatch from the time it was forward'ed until the corresponding backward.

- Gradient checkpointing can be used to trade extra computation for reduced activation memory: instead of storing all intermediate activations, some are recomputed during backward.



3. BatchNorm:

需要统计所有批次的均值与方差，但LLM中一般不用BN


- Standard BatchNorm layers (which compute stats over the full minibatch) break the microbatch independence assumption.

4. Communication-Compute Overlap:

通信与计算重叠

- Classic GPipe does not really overlap communication (transfers between GPUs) with compute, because each stage must finish its chunk before passing it!

假设batch分为两个micro_batch, gpu先接收到m1,后才能计算m1的forward,然后发送给gpu2, gpu2再进行forward



-